

A Bidirectional Neuro-Symbolic Foundation Pipeline for Continuous Sign Language AI

Closed-Loop Spatial Scaffolding, Causal Memory, and Anatomy-Aware Diffusion

Iniyan Joseph · Applicant, Doctor of Science in Technology (D.Sc.)

Department of Computer Science, School of Science, Aalto University, Espoo, Finland

Strategic Focus Alignment: ELLIS Finland Priority: *AI for Health & Assistive Digital Twins* · **Deadline:** Sept 21, 2026

Full Application Dossier: [Research Portal](#) | **Interactive Kinematics:** [WebGL Demo](#) | **LinkedIn:** [iniyandrews](#)

Executive Abstract & Project Vision

Continuous Sign Language Processing—spanning forward Sign Language Production (SLP: Text-to-Sign) and backward Sign Language Translation (SLT: Sign-to-Text)—represents a fundamental frontier in multimodal artificial intelligence and biomechanical digital twin modeling. Current computational paradigms exhibit severe structural limitations. In the forward pass, systems oscillate between rigid procedural avatars that lack natural biological momentum and unconstrained deep generative models (e.g., Pose-VQVAE, discrete diffusion) that suffer from quantization blur, spatial drift, and the smearing of centimeter-scale phonological minimal pairs. In the backward pass, streaming neural translators suffer from **Spatial State Collapse**—wherein sliding temporal context windows purge historical 3D spatial reference loci (R -loci), breaking directional verb agreement—and **Environmental Grounding Blindness**, failing to resolve polysemic depicting classifiers (such as handshake CL: C) without visual scene grounding. Crucially, current literature treats production and translation as disconnected unidirectional tasks.

This doctoral project establishes a unified, closed-loop **Bidirectional Neuro-Symbolic Foundation Model** for continuous sign language AI. In the **Forward Pass (Text-to-Sign)**, we decouple symbolic syntax from continuous physical naturalization across four specialized stages: (1) an SVO-to-Time-Topic-Comment (TTC) transfer grammar maintaining a dynamic 3D Spatial Registry with generative LLM procedural fallback for out-of-vocabulary terms; (2) a modular sub-lexical procedural engine compiling phonological primitives into a deterministic 52-joint SMPL-X skeletal scaffold ($\hat{\mathbf{P}}$) via closed-form 2-bone arm Inverse Kinematics (IK), 15-DOF 3-phalanx hand anatomy, Kapandji limits, and a quartic Biomechanical Strain/Pain Objective (P_{strain}); (3) a Continuous Residual Diffusion Transformer ($\mathbf{X}_{\text{final}} = \Pi_{\mathcal{M}}(\hat{\mathbf{P}} + \Delta\mathbf{X}_{\theta})$) projected onto the kinematic manifold to guarantee bone-length invariance and protect millimetric contact bounds ($r_{\text{tol}} \leq 5 \text{ mm}$); and (4) Adaptive Layer Normalization (AdaLN) for persona prosody and cross-lingual transfer. In the **Backward Pass (Sign-to-Text)**, we resolve streaming latency and environmental blindness via: (1) an anatomical perceptual front-end recovering 15-DOF hands and FLAME blendshapes in an invariant sternum frame; (2) continuous 3D loci raycasting from palm aiming and ocular gaze into an external, stateful **Causal Spatial Memory Bank** (M_t) decoupled in $\mathcal{O}(1)$ space from the LLM context window with deictic parity inversion ($\mathbf{P}_{\text{deictic}}$); and (3) an autoregressive decoder integrating *SceneGenie* 3D semantic scene graphs for real-world object grounding. Both modalities are unified by the **Spatial-SignLLM**, an autoregressive foundation model defined over a joint vocabulary $\mathcal{V}_{\text{total}} = \mathcal{V}_{\text{text}} \cup \mathcal{V}_{\text{phono}} \cup \mathcal{V}_{\text{spatial}}$ optimized via standard causal language modeling. We outline a phased 4-year plan under Aalto's **3-Paper Cumulative Strategy** (CVPR $\rightarrow$ ACL $\rightarrow$ NeurIPS) leveraging CSC's **EuroHPC LUMI supercomputer** and participatory co-design with the Finnish Association of the Deaf (*Kuurojen Liitto ry*).

1 Introduction, Societal Motivation & Global Human Impact

Over **70 million Deaf and hard-of-hearing individuals** worldwide communicate primarily through natural sign languages [12, 42]. Despite transformative advances in natural language processing (NLP) and large multimodal models (LMMs) for spoken communication, automated sign language technologies remain severely underdeveloped. Deaf citizens face persistent barriers in healthcare consultations, civic institutions, higher education, and digital services. Within the European Union, this gap conflicts with the European Accessibility Act (Directive 2019/882) and national legislation—most notably Finland's Sign Language Act (*Viittomakielilaki 352/2015*), which codifies the statutory right of signers to access public services in their mother tongue.

The core computational challenge of sign language stems from its physical medium. Sign language is not a 1D acoustic waveform or a sequence of linear text tokens; it is a **multi-channel, asynchronous, 4D spatio-temporal broadcast** operating across the 3D volume surrounding the human body over continuous time [12, 25, 42]. An authentic signing interaction coordinates four simultaneous visual channels:

- 1. Manual Articulations:** Rapid finger configurations (15 rotational degrees of freedom per hand), palm orientations, and continuous wrist trajectories in 3D Euclidean space.

2. **Non-Manual Signals (NMS):** Sub-pixel facial Action Units (eyebrow elevation AU1+AU2, furrowing AU4, mouth morphemes, ocular gaze vectors) that convey grammatical prosody, negation, and clause boundaries [32,33].
3. **Torso and Head Posture:** Spatial body shifts and head tilts marking discourse perspective and role shifting [27].
4. **3D Referential Space (*R-loci*):** The programmatic projection of invisible 3D coordinate anchors into the air around the torso to represent discourse entities, which directional agreement verbs physically interpolate between to govern grammatical agency [28,34].

Automating continuous sign communication requires solving two coupled problems: **Forward Generative Production (SLP: Text-to-Sign)**, synthesizing continuous, biologically natural, and grammatically sound 3D anatomical skeletal streams from text; and **Backward Perceptual Translation (SLT: Sign-to-Text)**, ingesting continuous video, tracking 3D spatial discourse anchors under streaming latency constraints (< 300 ms), and decoding fluent text while grounding gestures to the physical environment.

Historically, these tasks have been treated in isolation with monolithic deep neural networks. Because high-fidelity 3D motion-capture sign datasets are scarce (~ 80 hours globally), generative models suffer from anatomical hallucinations and quantization blur that smear phonological minimal pairs [47]. Conversely, streaming translators suffer from **Spatial State Collapse**, wherein sliding temporal windows purge the spatial memory required to resolve directional pronouns [24,34]. Hosted at the **ELLIS Institute Finland** and **Aalto University** under the supervision of **Dr. Azade Farshad**, this 4-year doctoral project establishes a bidirectional neuro-symbolic pipeline that guarantees structural grammatical integrity through deterministic procedural scaffolding while mastering physical fluidity, causal spatial memory, and foundation modeling through continuous diffusion and the Spatial-SignLLM.

2 Linguistic Taxonomy & 3D Spatial Phonology

To satisfy the strict grammatical standards of native Deaf signers, computational representations must be grounded in formal sign language linguistics.

2.1 Sub-Lexical Parameter Composition & Minimal Pairs

In formal sign linguistics (Stokoe, 1960 [42]; Battison, 1978 [4]), every lexical sign is decomposed into five discrete phonological parameters:

$$S = \langle \mathbf{H}, \mathbf{L}, \mathbf{M}, \mathbf{O}, \mathbf{F}_{\text{NMS}} \rangle \quad (1)$$

where **H** represents Handshape (finger joint angles), **L** denotes 3D anatomical Location target, **M** is Movement trajectory primitive, $\mathbf{O} \in \text{SO}(3)$ is wrist Orientation, and $\mathbf{F}_{\text{NMS}} = \langle \mathbf{b}, \mathbf{m}, \mathbf{g}, \mathbf{h} \rangle$ encodes facial Action Units (brows **b**, mouth **m**, gaze **g**, head tilt **h**).

Signs differing in exactly one parameter are **phonological minimal pairs**, where centimeter-scale deviations invert semantic meaning:

- ▶ **Location Minimal Pairs:** In American Sign Language (ASL), MOTHER is articulated with an open-5 handshape contacting the chin, whereas FATHER contacts the forehead. A vertical deviation of mere centimeters inverts maternal to paternal kinship.
- ▶ **Movement Minimal Pairs:** FATHER involves a single contact tap at the forehead, whereas GRANDFATHER shares the identical handshape and location but executes a double-arc outward ballistic bounce, requiring high-frequency temporal acceleration tracking.
- ▶ **Orientation Minimal Pairs:** NAME involves two H-handshapes tapping transversely; altering palm orientation from inward to downward produces an ungrammatical form.

2.2 Non-Manual Signals & 3D Spatial Syntax (*R-loci*)

Non-manual signals act as the visual equivalent of spoken prosody and syntactic clause delimiters [32,33]: raised eyebrows (AU1+AU2) mark Yes/No questions and conditionals, furrowed eyebrows (AU4) mark Wh-questions, and ocular gaze ($\mathbf{g}_{\text{gaze}}$) raycasts toward empty spatial loci to establish pronominal reference before manual indexing [12,41].

As modeled by Pfau et al. (2018) [34] and Lillo-Martin & Meier (2011) [28], signers project coordinate anchors into the air around their body, termed **Referential Loci (*R-loci*)**. In our interaction volume:

- ▶ Locus₁ (Signer Sternum): $[0.00, 0.00, 0.00]$ m, designating first-person (“I”, “Me”).



(a) ASL FATHER (single tap)



(b) ASL GRANDFATHER (double-arc outward bounce)

Figure 1: **Phonological Minimal Pairs:** Identical handshape and forehead location require millimetric spatial precision ($r_{\text{tol}} \leq 5$ mm) and high-frequency temporal acceleration tracking to prevent semantic inversion between kinship relations.



Figure 2: **Non-Manual Signals (NMS) as Visual Prosody:** Multi-channel face tracking capturing sub-pixel eyebrow Action Units (AU1+AU2 vs. AU4 for conditionals vs. questions), 3D ocular gaze raycasting vectors ($\mathbf{g}_{\text{gaze}}$) locking onto referential loci, and mouth morphemes.

- Locus₂ (Addressee): [0.00, 0.60, 0.00] m, designating second-person (“You”).
- Locus_{3a}, Locus_{3b} (Third-Party Referents): Right hemisphere [+0.35, 0.45, +0.05] m and Left hemisphere [−0.35, 0.45, +0.05] m.

Directional verbs physically traverse these loci: **Regular verbs** (GIVE, VISIT) move from Subject to Object (Locus_{3a} → Locus₁ for “Mother visits me”); **Backward verbs** (INVITE, TAKE) invert physical trajectory relative to agency, moving from Goal to Agent (Locus_{3a} → Locus₁ for “I invite Mother”) [30, 34]; and **Orientation-only verbs** (LOOK-AT, WARN) aim palm or index vectors directly toward the active referent.

2.3 Empirical Regularities in Sign Lexicons (ASL-LEX 2.0)


Mentored by Dr. Marco Marini (Sapienza University of Rome) to formalize sub-lexical decomposition, our parameter space is calibrated against **ASL-LEX 2.0** (2,723 signs across 54 dimensions [7, 38]). Prior linguistic analyses reveal critical structural constraints:

1. **Handshape Distribution Skew:** In accordance with Battison’s phonological laws [4], the top 6 unmarked handshapes (B, A/S, 1, 5, 0, C) represent $\approx 52\%$ of the lexicon, and the top 12 cover $\approx 78\%$. Calibrating 12 anatomical finger configurations provides immediate coverage of conversational vocabulary.
2. **Spatial Location Clustering:** Signs cluster tightly: $\approx 40\%$ occur in neutral space before the chest, $\approx 34\%$ cluster around cranial/facial landmarks (chin, forehead, cheek), and the remainder occur on the trunk or base hand, motivating millimetric precision on cranial zones.
3. **Battison’s Symmetry Constraints:** Two-handed asymmetric signs restrict the passive base hand to unmarked resting shapes, dramatically simplifying multi-hand coordination.
4. **Movement Dynamics:** Approximately 62% of signs involve single continuous paths (verbs), while 38% feature repeated contact taps or oscillations (nouns), parameterized via a repetition multiplier.

3 Literature Review & Rigorous Scientific Gap Analysis

3.1 Generative Production (SLP): Procedural Scaffolding vs. Neural Diffusion

Generative SLP has oscillated between procedural correctness and perceptual realism. Early procedural animation systems (VisiCast, JASigning) converted phonetic scripts (HamNoSys/SiGML) into deterministic avatars [29]. While guaranteeing bone-length invariance and joint limits, they produced robotic motion lacking biological momentum and co-articulation. Recent neural approaches using autoregressive



(a) Plain Verbs (LIKE, PAM)

(b) Regular Verbs (VISIT, $3a \rightarrow 1$)

(c) Backward Verbs (INVITE)

Figure 3: **Syntactic 3D Spatial Loci and Agreement Trajectories** [34]: Plain verbs anchor to the body, regular agreement verbs traverse from Subject to Object locus, and backward verbs invert physical trajectory relative to semantic agency.

Transformers, GANs, or discrete latent diffusion (e.g., G2P-DDM [47], Neural Sign Actors [3]) mapped glosses to vector-quantized latent spaces (Pose-VQVAE). However, codebook quantization introduces spatial blur that smears millimeter-scale minimal pairs. Conversely, continuous unconstrained skeletal diffusion models (e.g., MDM [44]) diffuse in Euclidean space $\mathbb{R}^{J \times 3}$; lacking kinematic manifold projection, they stretch bone lengths (“rubber fingers”) and breach physiological Range-of-Motion (ROM) limits. Simulator-guided diffusion (e.g., PhysDiff [49], BioMoDiffuse [5]) integrates rigid-body physics (Isaac Gym / OpenSim) for gross whole-body locomotion, but rigid-body simulators cannot resolve high-frequency 20-DOF finger dexterity, Kapandji thumb opposition, or delicate sub-centimeter ($r_{\text{tol}} \leq 5 \text{ mm}$) cranial contact boundaries where rigid contact penalties cause numerical chattering. Similarly, while recent transition diffusion models like StgcDiff [22] apply spatial-temporal graph convolutions (ST-GCN) as soft neural conditioning, soft feature conditioning cannot mathematically guarantee hard bone-length conservation (ℓ_{ij}) or prevent anatomical dislocations under fast strokes. Appearance-level models (e.g., DiffSign [26]) render photorealistic 2D RGB pixels from 2D poses, bypassing 3D kinematic fidelity. Our solution is **Manifold-Projected Continuous Residual Diffusion** ($\mathbf{X} = \Pi_{\mathcal{M}}(\hat{\mathbf{P}} + \Delta \mathbf{X}_{\theta})$), anchoring diffusion to exact procedural coordinates while enforcing iterative holonomic manifold constraints.

3.2 Perceptual Translation (SLT) & Multimodal Sign Foundation Models

Current SOTA Sign Language Translation treats the mapping from video to text as an end-to-end sequence-to-sequence problem ($\mathbf{V}_{1:T} \rightarrow \mathbf{W}$). In conversational settings, existing models exhibit three fundamental structural failure modes:

1. **Spatial State Collapse (The Streaming Bottleneck):** As empirically quantified by **Ranum, Hadfield, & Bowden (2026)** [36], real-time streaming translation ($< 300 \text{ ms}$ latency) under sliding temporal window truncation purges established spatial loci, inducing up to **74% directional agreement failures** on referential verbs (see also Kermani et al. [24]). While recent prompting heuristics (e.g., SpaMo [40]) append static frame-level 2D/3D Euclidean distances to LLM prompts, static text prompting lacks a stateful, persistent 3D spatial register that survives across multi-turn discourse. We resolve this by introducing our decoupled $\mathcal{O}(1)$ **Causal Spatial Memory Bank** (M_t), which persistently tracks 3D referent coordinates outside the temporal attention window.
2. **Environmental Grounding Blindness:** Pure-skeletal translators [11, 48] are blind to external physical space. When a signer points to a room object or uses an ambiguous depicting classifier (e.g., CL:C, representing a cup, cylinder, or vase) [43], skeletons lack visual grounding to identify the referent.
3. **Unidirectional Disconnection in Foundation Sign LLMs:** Recent multimodal foundation efforts bifurcate into disconnected paradigms: text-to-pose generation models (e.g., SignLLM [13]) generate raw 2D/3D skeletal keypoints directly from text prompts without 3D spatial referential loci, causing token explosion ($> 3,000 \text{ values/s}$) and phonetic minimal-pair smearing; conversely, recent sign-to-text translation (SLT) foundation models (e.g., ViPo-MLLM [45], SignBind-LLM [39]) align visual-pose features to language models, yet remain strictly unidirectional recognition systems lacking explicit 3D spatial memory (M_t) or bidirectional closed-loop 3D kinematic synthesis.

4 Proposed Technical Methodology: The Closed-Loop Pipeline

4.1 Architectural Blueprint & Functional Decoupling

The fundamental insight of our architecture is **functional decoupling**:

Research Challenge	Current SOTA Paradigm	Structural Limitation
Quantization Blur	Pose-VQVAE + Discrete Diffusion (G2P-DDM [47])	Codebook quantization rounds continuous coordinates, smearing minimal pairs.
Unconstrained Drift	Unconstrained Motion Diffusion (MDM [44])	Euclidean diffusion without kinematic manifold stretches bones (“rubber fingers”).
Locomotion vs. Dexterity	Simulator-Guided Diffusion (PhysDiff [49], BioMoDiffuse [5])	Rigid simulators model gross locomotion, failing on 20-DOF dexterity and cranial contacts.
Soft Graph Priors	Graph-Conditioned Diffusion (StgcDiff [22])	Soft ST-GCN conditioning cannot mathematically guarantee bone-length invariance (ℓ_{ij}).
Spatial State Collapse	Sliding-Window / Heuristics (Ranum [36], SpaMo [40])	Truncating context windows or static prompting purges historical R -loci across turns.
Environmental Blindness	Pure-Skeletal Translators (STMC [48])	Skeletons lack environmental context to resolve depicting classifiers ($CL:C$).
Unidirectional Disconnection	Text-to-Pose or SLT-only LLMs (SignLLM [13], ViPo-MLLM [45])	Lack 3D loci memory (M_t), 20-DOF kinematics, and bidirectional closed-loop synthesis.

Table 1: Core Algorithmic Challenges and SOTA Failure Modes in Sign Language AI.

- **Symbolic & Kinematic Algorithms** guarantee hard physical laws, anatomical joint limits, collision avoidance, and grammatical spatial indexing via **deterministic procedural compilers and closed-form analytical IK**.
- **Continuous Generative Neural Models** synthesize biological fluidity, micro-coarticulation, stylistic prosody, and commonsense reasoning via **continuous residual diffusion models and the BiSign-LLM**.

As shown in Figure 4, the pipeline couples the Forward Production Pass (SLP), the Backward Perceptual Pass (SLT), and the BiSign-LLM Multimodal Bridge into a closed-loop system.




Figure 4: The Unified Bidirectional Spatial Foundation Pipeline with the BiSign-LLM Multimodal Bridge.

4.2 The Forward Pass: Text-to-Sign Language Production (SLP)

The Forward Pass transforms spoken text into a continuous 60 Hz 3D avatar stream across four specialized stages:

4.2.1 Stage 1: Linguistic Parsing, 3D Spatial Registry & Generative LLM OOV Fallback

Natural sign languages function as Topic-Prominent languages governed by **Time-Topic-Comment (TTC)** structure [18, 27, 32]. Feeding linear spoken Subject-Verb-Object (SVO) text directly to an avatar produces ungrammatical pidgin. Stage 1 executes three operations:

1. **Syntactic Transfer Grammar (SVO $\rightarrow$ TTC):** Extracts temporal adverbials and fronts them to clause-initial position to establish the discourse timeline, deletes non-signing copulas (*is, are, was*), and shifts interrogative markers clause-final where facial markers (AU4) frame the proposition.
2. **Coreference Resolution & 3D Spatial Registry:** Allocates coordinate anchors in the signer's interaction volume: $\text{Locus}_1 = [0.00, 0.00, 0.00]$ m (Self), $\text{Locus}_{3a} = [+0.35, 0.45, +0.05]$ m (Primary referent), and $\text{Locus}_{3b} = [-0.35, 0.45, +0.05]$ m (Secondary referent). Directional verbs bind trajectory end-points directly between these coordinates. The sentence is compiled into a 3D spatial scene graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ encoded via a Relational Graph Transformer (RGT).
3. **Generative LLM Out-of-Vocabulary (OOV) Procedural Fallback:** Rather than freezing or falling back to tedious 8-second letter fingerspelling when encountering unseen terms [9, 31], we leverage the spatial commonsense reasoning of LLMs to synthesize **Depicting Signs (Classifier Predicates)** [12, 43] from atomic sub-lexical parameters. Crucially, the LLM does not emit continuous joint Euler angles directly; rather, it outputs high-level discrete phonological tokens—handshape primitives (**H**), movement paths (**M**), palm orientations (**O**), and contact locations (**L**)—which are subsequently validated against phonological collision rules and compiled into continuous kinematics by Stage 2.

Listing 1: Validated Pydantic 4D Phonetic JSON Interface with Procedural OOV Synthesis.

```
{
  "sentence_id": "AALTO_SVK_001", "asl_gloss": "YESTERDAY MOTHER 3a-VISIT-1",
  "sentence_nms": {"eyebrows": "NEUTRAL", "negation": false},
  "spatial_registry": {
    "Locus_3a": {"entity": "MOTHER", "coord": [0.35, 0.45, 0.05]},
    "Locus_1": {"entity": "SELF", "coord": [0.00, 0.00, 0.00]}
  },
  "tokens": [
    {"lemma": "YESTERDAY", "handshape": "HS_A_THUMB", "loc": [0.15, 0.10, 0.30], "mov": "CHEEK_TAP"},
    {"lemma": "MOTHER", "handshape": "HS_5_OPEN", "loc": [0.00, 0.08, 0.20], "mov": "CHIN_TAP"},
    {"lemma": "VISIT", "handshape": "HS_V_ALT", "start": [0.35, 0.45, 0.05], "end": [0.00, 0.00, 0.00]},
    {"lemma": "CAKEPOP", "oov_fallback": true, "procedural_tokens": {
      "idx": "CURVED", "mid": "CLOSED", "rng": "CLOSED", "pnk": "EXTENDED_SPREAD_R", "thm": "THM_PINCH_INDEX"}
    }
  ]
}
```

4.2.2 Stage 2: Modular Sub-Lexical Procedural Scaffolding & Biomechanical Muscle Engine

Stage 2 compiles the 4D JSON stream into a deterministic 52-joint SMPL-X skeletal scaffold $\hat{\mathbf{P}} \in \mathbb{R}^{T \times 52 \times 3}$ via analytical closed-form IK in $\mathcal{O}(1)$ time, guaranteeing anatomical feasibility before neural processing:

1. **15-Joint, 20-DOF Anatomical Hand Model & Musculature:** Fingers are modeled as authentic 3-phalanx kinematic chains across 15 articulated phalangeal joints: 4 fingers with Metacarpophalangeal (MCP: 2-DOF flex $[-15^\circ, 90^\circ]$, abd $[-20^\circ, 20^\circ]$), Proximal Interphalangeal (PIP: 1-DOF flex $[0^\circ, 100^\circ]$), and Distal Interphalangeal (DIP: 1-DOF flex $[0^\circ, 80^\circ]$) joints ($4 \times 4 = 16$ DOFs); and the thumb articulating across Trapeziometacarpal/CMC (2-DOF flex/opposition), MCP (1-DOF), and Interphalangeal/IP (1-DOF) joints (4 DOFs), providing 20 active mechanical degrees of freedom alongside modeled thenar and hypothenar muscular pads.
2. **Quartic Articular Range-of-Motion (ROM) Barrier Objective $P_{\text{strain}}(\theta)$:** Human motor control minimizes joint strain. We formalize this biological principle via an explicit quartic pain/strain penalty:

$$P_{\text{strain}}(\theta) = \sum_{j=1}^J \left(\frac{\theta_j - \theta_{j,\text{rest}}}{\theta_{j,\text{max}} - \theta_{j,\text{min}}} \right)^4 \quad (2)$$

Escalating sharply near physiological limits, resolving motion along the gradient of minimal strain naturally chooses the path of least physical resistance, preventing awkward joint hyperextensions.

While the skeletal scaffold $\hat{\mathbf{P}}$ is initialized via analytical closed-form IK, $P_{\text{strain}}(\theta)$ serves as a smooth non-linear interior barrier function defining the allowable kinematic manifold $\mathcal{M}$, evaluated during iterative manifold projection $\Pi_{\mathcal{M}}$ and offline trajectory optimization. Dynamic trajectory smoothness and minimum-jerk optimization ($\int \|\ddot{\mathbf{x}}\|^2 dt$) [19] are governed in continuous time by the residual diffusion score network.

3. **Multi-Digit Coordination & Soft-Body Collision Repulsion:** When distance between phalangeal segments falls below combined radii ($r_A + r_B$), an elastic repulsion vector prevents self-intersection. Dynamic priority scheduling orders digit flexion in folded signs (e.g., in M, N, T, pinky and ring curl before the thumb crosses and index folds over).
4. **Analytical Arm Kinematics with 2-DOF Clavicle & Humeral Swivel Orbit:** Given shoulder root $\mathbf{x}_{\text{shoulder}}$ and wrist $\mathbf{x}_{\text{wrist}}(t)$, the shoulder anchor is first modulated by a 2-DOF clavicle articulation model (15° – 20° elevation and protraction) to prevent reach singularities at cranial minimal-pair targets (e.g., FATHER at the forehead). Let $\vec{d} = \mathbf{x}_{\text{wrist}} - \mathbf{x}_{\text{clavicle}}$ ($d = \|\vec{d}\|$), with upper arm length L_1 and forearm length L_2 . The elbow flex angle θ_{elbow} is resolved in closed form via the Law of Cosines:

$$\theta_{\text{elbow}} = \arccos \left(\text{clamp} \left(\frac{d^2 - L_1^2 - L_2^2}{2L_1L_2}, -1, 1 \right) \right) \quad (3)$$

Humeral swivel redundancy $\phi \in [0, 2\pi)$ parameterizes elbow position along an orbit circle orthogonal to $\vec{u} = \vec{d}/d$, referenced to a biological resting plane with singularity damping at full arm extension.

5. **Facial Blendshape Binding ($\mathbf{F}_{\text{NMS}}$):** Non-manual clause modifiers are mapped into 50 FLAME blendshapes synchronized with the manual stroke phase.

4.2.3 Stage 3: Manifold-Constrained Kinematic Residual Diffusion

Pure procedural animation is structurally correct but mechanically stiff. Stage 3 employs continuous diffusion solely to synthesize biological momentum, micro-coarticulation, and natural fluidity atop the rigid scaffold:

$$\mathbf{X}_{\text{final}} = \Pi_{\mathcal{M}}(\hat{\mathbf{P}} + \Delta\mathbf{X}_{\theta}) \quad (4)$$

where $\Delta\mathbf{X}_{\theta}$ is the predicted residual micro-displacement. Grounded in **Dr. Farshad's methodologies in topologically constrained generative models** (SCOPE [16], SceneGenie [14]), we enforce Riemannian manifold projection ($\Pi_{\mathcal{M}}$) via iterative holonomic constraint projection using the SHAKE/RATTLE molecular dynamics algorithm [2, 37] (converging in ≤ 3 – 5 iterations) adapted to skeletal kinematic trees:

1. **Bone-Length Invariance ($\mathcal{L}_{\text{graph}}$):** Enforces exact distance constraints ($\|\mathbf{x}_i(t) - \mathbf{x}_j(t)\|_2 = \ell_{ij}$) across all 52 skeletal connections and 15 phalangeal segments, permanently eliminating bone stretching.
2. **Character-Adaptive Contact Tolerance (r_{tol}):** Enforces tight millimetric bounds ($r_{\text{tol}} \leq 5$ mm) at critical contact landmarks (e.g., chin vs. forehead for MOTHER vs. FATHER), while permitting natural compliance (20–30 mm) during unconstrained movement in neutral space.
3. **Pain Envelope Preservation:** Ensures neural residual updates respect the joint limit and strain boundaries established by $P_{\text{strain}}(\theta)$.

The denoiser is pretrained on **Motion-X++** (19.5M frames) [50] (regularized by $\Pi_{\mathcal{M}}$), then fine-tuned on How2Sign's 3D Panoptic subset (~ 3 h mocap) [11].

4.2.4 Stage 4: Stylistic Prosody & Universal Cross-Lingual Transfer

A style vector $\mathbf{s} \in \mathbb{R}^{512}$ modulates the denoiser via Adaptive Layer Normalization (AdaLN) to synthesize distinct signer personas, signing tempos, and emotional prosody. Swapping sub-lexical dictionary tables enables zero-shot cross-lingual transfer from ASL to Finnish Sign Language (SVK) and other sign languages with a detailed enough sub-lexical dictionary.

4.3 The Backward Pass: Sign-to-Text Language Translation (SLT)

The Backward Pass perceives and translates continuous signing video back into spoken text through three interpretable stages:

4.3.1 Stage 1: Multi-Channel Perceptual Front-End & Domain-Optimized Tracking

Operating over dedicated signers facing the camera, the perceptual front-end recovers 3D anatomical articulators:

1. **20-DOF Hand Mesh Reconstruction:** High-dexterity tracking via HaMeR [51] extracts rotational joint angles across all 20 active mechanical degrees of freedom (15 joints):

$$\mathbf{q}_{\text{hand}} = [\mathbf{q}_{\text{thumb}} \in \mathbb{R}^4, \mathbf{q}_{\text{index}} \in \mathbb{R}^4, \mathbf{q}_{\text{middle}} \in \mathbb{R}^4, \mathbf{q}_{\text{ring}} \in \mathbb{R}^4, \mathbf{q}_{\text{pinky}} \in \mathbb{R}^4]^T \in \mathbb{R}^{20} \quad (5)$$

2. **Facial Action Units & Gaze Vector:** Tracks 50 FLAME blendshapes (eyebrows, mouth morphemes) and the 3D ocular gaze unit vector $\mathbf{g}_{\text{gaze}}(t) \in \mathbb{R}^3$.
3. **Body-Centric Sternum Normalization Frame:** Skeletal keypoints $\mathbf{x}_{\text{world}} \in \mathbb{R}^3$ are normalized to an invariant coordinate frame centered at the signer's sternum $([0, 0, 0])$ and scaled by biacromial shoulder width (W_{shoulder}), eliminating perspective and scale variances.

4.3.2 Stage 2: 3D Spatial Loci Raycasting & Central Memory Bank (M_t)

Stage 2 tracks spatial interactions and maintains discourse memory across turns:

- ▶ **Palm & Gaze Aiming Raycasting:** Continuous 3D aiming rays are cast from the normalized palm orientation vector ($\vec{v}_{\text{aim}} = \mathbf{R}_{\text{wrist}} \cdot \mathbf{u}_{\text{point}}$) and ocular gaze vector $\mathbf{g}_{\text{gaze}}(t)$, intersecting active 3D spatial loci in the interaction volume.
- ▶ **Decoupled Central Memory Bank (M_t):** Active referents ($K \leq 8$ entities) are stored in an external stateful memory bank:

$$M_t = \{\mathbf{m}_k = \langle \mathbf{x}_k, \mathbf{e}_k, \alpha_k, \tau_k \rangle\}_{k=1}^K \quad (6)$$

where $\mathbf{x}_k \in \mathbb{R}^3$ is the physical coordinate anchor, $\mathbf{e}_k$ is the entity embedding, α_k is the salience weight, and τ_k is the last interaction timestamp. Decoupling M_t into an $\mathcal{O}(1)$ external register allows entities to persist across long dialogues without inflating the LLM context or causing spatial state collapse when sliding temporal windows advance. Crucially, while articulatory tracking is normalized to the sternum for scale invariance, spatial referents $\mathbf{x}_k$ in M_t are parameterized in an egocentric horizontal ground-plane frame decoupled from torso yaw, preventing discourse anchors from erroneously rotating when a signer executes a physical role-shift (15° – 30° torso turn).

- ▶ **Deictic Frame-of-Reference Inversion ($\mathbf{R}_y(\pi)$):** In face-to-face dialogue, the coordinate frames of signer and avatar face each other along the line of sight. To preserve biological body and hand chirality while aligning spatial reference loci during turn-taking, we apply a proper rigid $SO(3)$ deictic rotation with conversational offset $\mathbf{t}_{\text{conv}}$:

$$\mathbf{x}_{\text{avatar}} = \mathbf{R}_y(\pi) \cdot \mathbf{x}_{\text{signer}} + \mathbf{t}_{\text{conv}}, \quad \mathbf{R}_y(\pi) = \text{diag}(-1, 1, -1) \quad (7)$$

Crucially, because $\det(\mathbf{R}_y(\pi)) = (-1)(1)(-1) = +1 \in SO(3)$, this transformation preserves physical chirality (a right hand remains an authentic right hand), while correctly mapping the signer's right hemisphere (contralateral to the addressee's left visual field) to the corresponding shared deictic locus.

- ▶ **Spatial Hysteresis & Thematic θ -Role Matching:** To mitigate high-frequency monocular tracking noise without sacrificing precision, aiming rays are temporally integrated across a 5–10 frame window (~ 160 ms). Instead of a naive rigid threshold, intersection with active referents in M_t is resolved via an adaptive probabilistic Gaussian tolerance kernel $\mathcal{N}(\mathbf{x}_k, \Sigma_{k,t})$, where spatial covariance $\Sigma_{k,t}$ is updated dynamically by an Extended Kalman Filter (EKF). Furthermore, phonological contrast in natural sign discourse ensures that distinct entities inhabit well-separated canonical signing sectors (contralateral, center, ipsilateral) with angular separation $> 30^\circ$, reliably exceeding the monocular tracking uncertainty cone. Directional verbs are resolved across three phonological classes [30]: Path verbs (GIVE) via trajectory start/end ($\mathbf{p}_{\text{start}} \rightarrow \mathbf{p}_{\text{end}}$); Backward verbs (INVITE) via lexical mapping inverting Goal to Source; and Orientation verbs (LOOK-AT) via $\vec{v}_{\text{aim}}$.

4.3.3 Stage 3: LLM Translation, Memory Management & Scene-Graph Grounding

Stage 3 decodes fluent text while governing memory and grounding gestures in the physical environment:

- ▶ **Active Memory Governance:** The LLM manages M_t dynamically: registering new loci $\mathbf{x}_{\text{new}}$ upon introduction, refreshing salience α_k upon re-reference, and evicting completed topics via soft exponential decay $\alpha_k(t)$ in $\mathcal{O}(1)$ time.
- ▶ **Environmental Grounding via Semantic Scene Graphs (SceneGenie [14]):** When a signer points to a real-world object (e.g., a coffee cup, door, or medical tool) or uses an ambiguous polysemic depicting classifier (CL:C), skeletal tracking cannot disambiguate the referent. We structure the room environment into a 3D semantic scene graph $\mathcal{G}_{\text{env}} = (\mathcal{V}_{\text{obj}}, \mathcal{E}_{\text{rel}})$. Intersecting the pointing ray $\vec{v}_{\text{aim}}$

with object bounding volumes allows the LLM to instantly resolve real-world referents without heavy raw video compute.

4.4 The BiSign-LLM: Unified 3D Spatial-Lexical Foundation Modeling

Bidirectional synergy between continuous physical kinematics and symbolic linguistics is unified through an autoregressive **Bidirectional Spatial Sign Large Model (BiSign-LLM)**. Rather than attempting to train high-variance reinforcement learning policies or unstable GAN discriminators across the continuous 52-joint boundary, we treat spatial sign communication as a **hierarchical multimodal language modeling task**.

Discrete 3D Coordinate Tokenization: Drawing on embodied robotics and vision-language-action models (RT-2 [6], PaLM-E [10], SpatialVLM [8]), large Transformer backbones natively ingest and emit physical spatial coordinates when continuous 3D space is discretized into coordinate tokens. While prior architectures such as SignLLM [13] pioneered text-to-pose generation, they autoregressively emit continuous joint coordinates directly without 3D spatial referential memory, leading to severe token explosion ($> 3,000$ values/s) and minimal-pair blurring. In contrast, our **BiSign-LLM** unifies spatial linguistics and biomechanics by defining a multimodal spatial-lexical vocabulary:

$$\mathcal{V}_{\text{total}} = \mathcal{V}_{\text{text}} \cup \mathcal{V}_{\text{phono}} \cup \mathcal{V}_{\text{spatial}} \quad (8)$$

where $\mathcal{V}_{\text{text}}$ is the standard subword vocabulary for spoken text, $\mathcal{V}_{\text{phono}}$ is the sub-lexical parameter alphabet $\mathcal{S}$ representing discrete handshapes and movement primitives, and $\mathcal{V}_{\text{spatial}}$ denotes discrete coordinate tokens partitioning the canonical interaction volume $[-1, 1]^3$ into spatial bins:

$$\mathbf{z}_{\text{locus}} = \text{Embed}_{\text{spatial}}(\langle x, y, z \rangle) \in \mathbb{R}^{d_{\text{model}}} \quad (9)$$

The entire foundation model is trained end-to-end via standard causal autoregressive language modeling cross-entropy over discrete token sequences:

$$\mathcal{L}_{\text{LM}} = - \sum_{i=1}^N \log P(w_i | w_{<i}), \quad w_i \in \mathcal{V}_{\text{total}} \quad (10)$$

How It Operates in the Bidirectional Loop:

1. **Forward Direction (Macro-Compiler at 1–2 Hz):** Direct autoregressive prediction of 60 Hz joint torques causes catastrophic token explosion ($60 \text{ frames/s} \times 52 \text{ joints} = 3,120 \text{ values/s}$). The BiSign-LLM resolves this by operating exclusively at the **discourse rate (1–2 Hz)**, emitting structured token sequences per concept:

[TIME: YESTERDAY] $\rightarrow$ [LOC_3a: MOTHER, (+0.35, 0.45, +0.05)] $\rightarrow$ [VERB: VISIT, 3a $\rightarrow$ 1]

When encountering unseen terms (e.g., 'CakePop'), the model emits discrete procedural depicting tokens (PINCH_INDEX, pinky spread) rather than continuous joint angles. This discrete 1–2 Hz blueprint parameterizes Stage 2 (closed-form 20-DOF procedural IK) and Stage 3 (residual diffusion), which synthesize smooth 60 Hz physics. The LLM acts as the “planner” while the kinematic pipeline acts as the “motor execution system.”

2. **Backward Direction (Sign-to-Text Decoding):** In reverse, our video perceptual front-end and memory bank M_t extract 20-DOF hand configurations $\mathcal{S}_t$ and tracked 3D coordinate tokens $\mathcal{V}_{\text{spatial}}$ from video. The BiSign-LLM ingests these multimodal tokens; self-attention heads bind the spatial trajectory endpoints back to the entities in M_t , correctly translating directional agreement verbs (VISIT, INVITE) into fluent spoken sentences without context window collapse.
3. **Optimization Stability on EuroHPC LUMI:** Pretraining across How2Sign’s Panoptic subset and Motion-X++ (19.5M video-estimated frames regularized via $\Pi_{\mathcal{M}}$) with pseudo-grounded scaffolds extracted by our perception front-end uses standard cross-entropy, eliminating the high gradient variance of continuous policy gradients.

4.5 Empirical De-risking: Working 20-DOF Kinematic Prototype

To de-risk the technical methodology against concerns of doctoral tractability and demonstrate that the proposed neuro-symbolic pipeline is practically realizable, we have independently developed and

deployed an interactive 20-DOF WebGL kinematics engine ([anatomical_hand.html](#)). Engineered deliberately as a lightweight, zero-dependency client-side prototype capable of executing smoothly across any consumer browser or mobile device, the engine compiles symbolic sub-lexical instructions into deterministic 3D hand configurations in closed-form $\mathcal{O}(1)$ time under anatomical joint limits and real-time tendon spring kinetics. While this interactive demonstrator proves topological feasibility on everyday hardware, the full non-convex quartic strain energy landscape $P_{\text{strain}}(\theta)$ and multi-body capsule collision solvers are evaluated in offline batch optimization (PyTorch/Taichi) for large-scale training on EuroHPC LUMI.

Crucially, to validate the tractability of generative OOV depiction in practice, when prompted with the novel out-of-vocabulary concept 'CakePop' (a spherical confection atop a stick), the foundation LLM does not hallucinate raw continuous joint angles. Instead, it operates strictly at the discrete phonological level—reasoning over the object's physical affordances to emit the discrete sub-lexical tokens:

$$\text{Tokens}(\text{CakePop}) = \{\text{idx} : \text{CURVED}, \text{mid} : \text{CLOSED}, \text{rng} : \text{CLOSED}, \\ \text{pnk} : \text{EXTENDED_SPREAD_R}, \text{thm} : \text{THM_PINCH_INDEX}\} \quad (11)$$

which are then deterministically compiled by the kinematic engine into anatomically valid 3D joint configurations. While ASL-LEX 2.0 serves as a closed lexical database of 2,723 established signs to calibrate canonical phonological primitive distributions and joint limit filters, open-domain productive depicting classifiers are validated against formal classifier grammar [43] and evaluated via Deaf expert panel rating (assessing naturalness, semantic plausibility, and physical affordance fidelity). This working proof-of-concept establishes that the candidate already commands the foundational 3D graphics, analytical kinematics, and neuro-symbolic compilation infrastructure required to execute Track A and Track C without speculative startup delays.

5 Methodological Synergy with Target Lab at Aalto University & ELLIS Finland

Continuous sign language processing represents a high-impact application domain for **Dr. Azade Farshad's** research group at **Aalto University / ELLIS Institute Finland**. As established in Section 4, our bidirectional neuro-symbolic framework directly integrates and extends her group's pioneering methodologies in geometric deep learning and generative modeling:

1. **Scene-Graph Generative Conditioning (Extending *SceneGenie* [14]):** In *SceneGenie*, Dr. Farshad demonstrated that conditioning diffusion models on structured semantic scene graphs significantly improves visual coherence. Sign discourse syntax operates precisely as a 3D spatial scene graph $\mathcal{G} = (\mathcal{V}_{\text{entities}}, \mathcal{E}_{\text{relations}})$, where nodes encode coordinate loci (L_1, L_{3a}, L_{3b}) and edges encode directional agreement predicates. Encoded via a Relational Graph Transformer (RGT), these embeddings condition the residual diffusion denoiser, eliminating minimal-pair drift. In reverse, scene graphs ground physical room objects to disambiguate pointing gestures.
2. **Topologically Constrained Generative Modeling (Extending *SCOPE* [16] & *HieraSurg* [15]):** In surgical video generation and anatomical segmentation, Dr. Farshad established that biological structures require explicit graph-based continuity regularizers. High-DOF articulated hands and arms exhibit identical topological vulnerabilities during neural diffusion. Adapting these graph constraints guarantees bone-length invariance ($\mathcal{L}_{\text{graph}}$) and biological strain limits ($\mathcal{P}_{\text{pain}}$), preventing joint dislocations.
3. **ELLIS Priority "AI for Health & Society":** According to the WHO, health encompasses complete physical, mental, and social well-being. Chronic communication isolation represents a severe social determinant of health. This bidirectional system acts as an **Assistive Communication Digital Twin**, dismantling systemic exclusion across healthcare, public services, and daily life.

6 Multi-Tier Evaluation Framework

To benchmark generative biomechanical accuracy (SLP) and translation fidelity (SLT), we establish the multi-tier framework summarized in Table 2. Crucially, existing continuous sign datasets (How2Sign [11], PHOENIX-2014T) contain text transcripts and keypoints, but lack 3D spatial referent ground truth (no labeled 3D R -loci or co-reference chains). Rather than conditioning doctoral milestones on massive manual re-annotation, our primary validation uses **Task-Grounded Downstream Evaluation** via

Argument Translation F1 (subject-object pronoun precision/recall in English/Finnish) and **Agreement-Gated BLEU-4**, isolating sentences strictly requiring spatial agreement. As a secondary direct check, if feasible within community partnership resources, we pilot a targeted **Spatial Agreement Benchmark (SAB-100)**, annotating 100 multi-turn clips with Deaf linguists from *Kuurojen Liitto ry* (scalable up to 500 clips if resources permit).

Dimension	Evaluation Metric & Formulation	Design Target	Comparative Literature Baseline
Landmark Precision	Anchor-Fidelity Error (EFE_{mm}) : Euclidean error on finger-tip/cranial keypoints.	≤ 15 mm	G2P-DDM [47] ($\approx 35\text{--}40$ mm)
Spatial Grounding	Perplexity & Spatial Locus Alignment (SLA_{mm}) : Causal LM perplexity and 3D coordinate precision.	PPL ≤ 8.5 , SLA ≤ 12 mm	Unconditioned LLMs / Disconnected tokenizers
Kinematic Realism	Fréchet Video Distance (FVD) and Kinematic Jerk ($\ \ddot{x}\ ^2$).	FVD ≤ 40.0	Human Motion Diffusion (MDM) [44]
Minimal-Pair Accuracy	Spatio-temporal classifier distinguishing minimal pairs (MOTHER vs. FATHER).	$\geq 92\%$ accuracy	Unconstrained generation baselines ($< 70\%$)
Agreement Resolution	Argument F1 & Direct ARA : Downstream Argument F1 on How2Sign; direct ARA on SAB pilot (if feasible).	Arg-F1 $\geq 88\%$, ARA $\geq 90\%$	Sliding-window SLT models [48] ($\approx 64\%$)
Continuous Translation	BLEU-4 & WER : Evaluated across continuous benchmarks (How2Sign, PHOENIX-2014T).	BLEU-4 ≥ 28.5 , WER ≤ 24.0	CorrNet+ [23], STMC [48]
Inference Latency	End-to-end streaming latency on modern GPU hardware.	< 100 ms (> 60 fps)	Dense Vision-Language Models (> 400 ms)
Everyday Chat MOS	Native Deaf User Evaluation: Likert scale (1–5) on Conversational Fluency and Spatial Clarity.	MOS $\geq 4.2/5.0$	Procedural animation avatars [29]

Table 2: Multi-Tier Evaluation Matrix balancing linguistic precision, foundation model perplexity, and physical realism.

7 4-Year Doctoral Research Work Plan: Parallel Tracks & Staggered Milestones

Adhering to Aalto University’s articles-based doctoral dissertation model (requiring 3–4 top peer-reviewed publications), the research progresses across two **Concurrent Development Tracks (Track A: Kinematic Synthesis; Track B: Perceptual Tracking)** that execute in parallel from Month 1, coupled via the shared sub-lexical vocabulary S and body-centric frame, converging into the foundation model (**Track C**):

- ▶ **Years 1–2: Parallel Foundations (Kinematic Synthesis & Streaming Perception)**
 - ▶ **WP1 (M1–M16, Track A: Kinematic Scaffolding & Residual Diffusion)**: Secure ethical approval from Aalto Research Ethics Committee; convene Deaf Advisory Board with *Kuurojen Liitto ry*. Implement procedural scaffolding with regularized 2-bone arm IK, seeded with **3D-LEX v1.0** [35] and **ASL-LEX 2.0** [38]. Deploy Continuous Residual Diffusion with localized tolerance bubbles ($r_{\text{tol}} \leq 5$ mm), graph invariance ($\mathcal{L}_{\text{graph}}$), and pain loss ($\mathcal{P}_{\text{pain}}$). Fine-tune on How2Sign’s 3D Panoptic subset [11].
 - ▶ **WP2 (M1–M16, Track B: Streaming Perception & Causal Spatial Memory)**: Developed concurrently with Track A. Deploy HaMeR hand tracking and sternum normalization. Build $\mathcal{O}(1)$ Causal Spatial Memory Bank (M_t) with soft exponential decay, gaze-pointing ray convergence, and θ -role matching resolving regular, backward (INVITE), and orientation verbs (LOOK-AT) via aiming vector $\vec{v}_{\text{aim}}$.
 - ▶ **Paper 1 Deliverable (Month 17, March 2028)**: *Anatomically Constrained Residual Diffusion for Articulated Sign Kinematics*. **Target Venue: ECCV 2028**.
- ▶ **Year 3: Track Convergence, BiSign-LLM Pretraining & Portability**
 - ▶ **WP3 (M17–M26, Track B/C: Dual-Rate Streaming)**: Deploy fast 60 Hz kinematics coupled with slow 1 Hz scene context (z_{scene}) and entropy-gated active visual interrupts (v_{patch}). Unify phonological primitives and 3D coordinate tokens into $\mathcal{V}_{\text{total}}$.
 - ▶ **WP4 (M24–M32, Track C: BiSign-LLM Pretraining on EuroHPC LUMI)**: Pretrain autoregressive BiSign-LLM on CSC’s **EuroHPC LUMI supercomputer** (AMD MI250X nodes) across How2Sign and Motion-X++ (19.5M frames) [50] using PyTorch FSDP and ROCm FlashAttention. Pretrain 1–2 Hz macro-compiler and depiction classifier synthesis for OOV terms.
 - ▶ **Paper 2 Deliverable (Month 33, June 2029)**: *BiSign-LLM: 3D Spatial-Lexical Foundation Modeling for Continuous Sign Language Translation*. **Target Venue: EMNLP 2029** (or ARR Spring cycle).

- ▶ *WP5 (M31–M35, Track C: Cross-Lingual Portability & Everyday Chat Benchmark)*: Adapt sub-lexical dictionaries from ASL to Finnish Sign Language (**SVK**). Couple BiSign-LLM with residual diffusion for closed-loop bidirectional translation. Evaluate on an **Everyday Conversational Benchmark** with native Deaf signers from *Kuurojen Liitto ry*.
- ▶ **Paper 3 Deliverable (Month 36, September 2029)**: *Multimodal Spatial Foundation Modeling for Open-Domain Assistive Communication*. **Target Venue: ICLR 2030**.
- ▶ **Year 4: Pre-Examination, Community Synthesis & Statutory Defense**
 - ▶ *WP6 (M37–M41, Oct 2029–Feb 2030)*: Comprehensive Deaf community user study with certified Deaf/SVK consultants ($MOS \geq 4.2/5.0$) and open-source foundation release.
 - ▶ *Statutory Pre-Examination (Esitarkastus, M42–M47, Mar–Aug 2030)*: Cumulative dissertation compilation and formal pre-examination by external opponents appointed by Aalto Doctoral Programme Committee in Science.
 - ▶ *Public Defense (Väitöstilaisuus, Month 48, September 2030)*: Formal doctoral defense at Aalto University.

8 Computational Feasibility, Ethical Governance & Risk Mitigation

High-Performance Computing Strategy (EuroHPC LUMI): Kinematic pretraining on Motion-X++ (19.5M video-estimated frames regularized via $\Pi_{\mathcal{M}}$) [50] and BiSign-LLM leverage Aalto’s direct access to the **EuroHPC LUMI consortium** (CSC, Jyväskylä, Finland). LUMI-G nodes feature **AMD Instinct MI250X GPUs**. Native **AMD ROCm / HIP**, FlashAttention, Composable Kernel, and PyTorch FSDP enable multi-node scaling. Datasets are hosted on LUMI’s parallel **Lustre filesystem** (/scratch) packaged into indexed **WebDataset** tar shards for high-throughput streaming.

Ethical Governance & Deaf Community Participatory Co-Design: To guarantee ethical integrity and avoid technological paternalism, formal ethical review is conducted by the **Aalto University Research Ethics Committee** prior to human evaluations. In partnership with *Kuurojen Liitto ry* (Finnish Association of the Deaf), an ongoing **Deaf Advisory Board** provides continuous feedback on blendshapes, cultural nuance, and linguistic naturalism from Month 1.

3-Tier Contingency Fallback Hierarchy: To guarantee timely completion adhering to Aalto’s 3-paper doctoral requirements:

- ▶ *Tier 1 (Core Non-Negotiable)*: Full ASL bidirectional procedural-diffusion pipeline, causal memory bank, mathematical loss formulations, and automated quantitative benchmarking (Years 1–3).
- ▶ *Tier 2 (Cross-Lingual Flexibility)*: If SVK extraction encounters delays, transfer focuses on calibrated ASL agreement benchmarks before scaling to open-domain SVK continuous discourse.
- ▶ *Tier 3 (User Study Scalability)*: If participant recruitment bottlenecks arise, evaluation transitions flexibly to a panel of 5 certified Deaf interpreters using standardized Likert rubrics.

References

-
- [1] Y. Aloimonos, I. Weiss, and A. Bandyopadhyay, “Active perception,” *Science*, vol. 238, no. 4831, pp. 1269–1275, 1988.
 - [2] H. C. Andersen, “RATTLE: A ‘velocity’ version of the SHAKE algorithm for molecular dynamics calculations,” *Journal of Computational Physics*, vol. 52, no. 1, pp. 24–34, 1983.
 - [3] V. Baltatzis, R. A. Potamias, E. Ververas, G. Sun, J. Deng, and S. Zafeiriou, “Neural Sign Actors: A diffusion model for 3D sign language production from text,” in *Proc. IEEE/CVF CVPR*, pp. 1985–1995, 2024.
 - [4] R. Battison, *Lexical Borrowing in American Sign Language*, Linstok Press, Silver Spring, MD, 1978.
 - [5] P. Peng et al., “BioMoDiffuse: Biomechanically compliant human motion diffusion model with musculoskeletal simulation,” in *Proc. ECCV*, 2024.
 - [6] A. Brohan, N. Brown, J. Carbajal, Y. Chebotar, X. Chen, et al., “RT-2: Vision-language-action models transfer web knowledge to robotic control,” in *Proc. 7th Conf. on Robot Learning (CoRL)*, 2023.
 - [7] N. K. Caselli, Z. S. Sehyr, A. M. Cohen-Goldberg, and K. Emmorey, “ASL-LEX: A lexical database for American Sign Language,” *Behavior Research Methods*, vol. 49, no. 2, pp. 784–801, 2017.
 - [8] B. Chen, Z. Xu, S. Kirmani, B. Ichter, D. Driess, P. Florence, D. Sadigh, et al., “SpatialVLM: Endowing vision-language models with spatial reasoning capabilities,” in *Proc. IEEE/CVF CVPR*, pp. 14455–14465, 2024.
 - [9] D. Cokely, *Interpretation: A Sociolinguistic Model*, Linstok Press, Burtonsville, MD, 1992.
 - [10] D. Driess, F. Xia, M. S. Sajjadi, C. Lynch, A. Chowdhery, et al., “PaLM-E: An embodied multimodal language model,” in *Proc. 40th ICML*, pp. 8469–8488, 2023.
 - [11] A. Duarte, S. Palaskar, L. Ventura, D. Ghadiyaram, K. DeHaan, F. Metze, J. Torres, and X. Giro-i-Nieto, “How2Sign: A large-scale multimodal dataset for continuous American Sign Language,” in *Proc. IEEE/CVF CVPR*, pp. 2735–2744, 2021.
 - [12] K. Emmorey, *Language, Cognition, and the Brain: Insights from Sign Language Research*, Lawrence Erlbaum Associates, Mahwah, NJ, 2002.
 - [13] E. Fang et al., “SignLLM: Sign languages production large language models,” *arXiv preprint arXiv:2405.10718*, 2024.
 - [14] A. Farshad et al., “SceneGenie: Scene graph-guided diffusion models for complex scene synthesis,” *arXiv preprint*, 2024.
 - [15] A. Farshad et al., “HieraSurg: Hierarchy-aware diffusion model for surgical video generation,” in *Proc. MICCAI*, 2025.
 - [16] A. Farshad et al., “SCOPE: Structural continuity preservation for medical image segmentation via graph constraints,” *arXiv preprint*, 2024.
 - [17] C. Feichtenhofer, H. Fan, J. Malik, and K. He, “SlowFast networks for video recognition,” in *Proc. IEEE/CVF ICCV*, pp. 6202–6211, 2019.
 - [18] S. D. Fischer, “Influences on word order change in American Sign Language,” in *Word Order and Word Order Change*, C. N. Li, Ed., University of Texas Press, pp. 1–25, 1975.
 - [19] T. Flash and N. Hogan, “The coordination of arm movements: an experimentally confirmed mathematical model,” *Journal of Neuroscience*, vol. 5, no. 7, pp. 1688–1703, 1985.
 - [20] A. Graves, “Adaptive computation time for recurrent neural networks,” *arXiv preprint arXiv:1603.08983*, 2016.
 - [21] D. He, Y. Xia, T. Qin, L. Wang, N. Yu, T. Liu, and W. Ma, “Dual learning for machine translation,” in *Proc. NeurIPS*, vol. 29, pp. 820–828, 2016.
 - [22] J. He, J. He, S. Tang, H. Ben, L. Cheng, and R. Hong, “StgcDiff: Spatial-temporal graph condition diffusion for sign language transition generation,” in *Proc. 3rd ACM MMGR at ACM Multimedia*, 2025.
 - [23] L. Hu, W. Feng, L. Gao, Z. Liu, and L. Wan, “CorrNet+: Sign language recognition and translation via spatial-temporal correlation,” *arXiv preprint arXiv:2404.11111*, 2024.
 - [24] A. Kermani, H. Irani, D. Ross, and V. Metsis, “Temporal-linguistic adaptive streaming for continuous sign language translation,” in *Proc. 4th Workshop on Advances in Language and Vision Research (ALVR) at ACL*, pp. 239–248, 2026.
 - [25] E. S. Klima and U. Bellugi, *The Signs of Language*, Harvard University Press, Cambridge, MA, 1979.
 - [26] S. Krishnamurthy, V. Bhat, and A. Jain, “DiffSign: AI-assisted generation of customizable sign language videos with enhanced realism,” in *Proc. ECCV Workshops*, 2024.
 - [27] S. K. Liddell, *Grammar, Gesture, and Meaning in American Sign Language*, Cambridge University Press, 2003.
 - [28] D. Lillo-Martin and R. P. Meier, “On the linguistic status of ‘agreement’ in sign languages,” *Theoretical Linguistics*, vol. 37, no. 3-4, pp. 95–141, 2011.
 - [29] J. McDonald, R. Wolfe, J. Schnepf, et al., “An automated technique for real-time production of lifelike animations of American Sign Language,” *Universal Access in the Information Society*, vol. 15, no. 4, pp. 551–566, 2016.
 - [30] I. Meir, “Syntactic-semantic interaction in Israeli Sign Language verbs: The case of backwards verbs,” *Sign Language & Linguistics*, vol. 1, no. 1, pp. 3–37, 1998.
 - [31] J. Napier, *Sign Language Interpreting: Linguistic Coping Strategies*, Douglas McLean Publishing, Bristol, 2002.

- [32] C. Neidle, J. Kegl, D. MacLaughlin, B. Bahan, and R. G. Lee, *The Syntax of American Sign Language: Functional Categories and Hierarchical Structure*, MIT Press, Cambridge, MA, 2000.
- [33] R. Pfau and J. Quer, “Nonmanuals: The expression of syntactic and prosodic structure,” in *Sign Language: An International Handbook*, De Gruyter Mouton, pp. 381–402, 2010.
- [34] R. Pfau, M. Salzmann, and M. Steinbach, “The syntax of sign language agreement: Common ingredients, but unusual recipe,” *Glossa: A Journal of General Linguistics*, vol. 3, no. 1, p. 107, 2018.
- [35] O. Ranum, G. Otterspeer, J. I. Andersen, R. G. Belleman, and F. Roelofsen, “3D-LEX v1.0 – 3D lexicons for American Sign Language and Sign Language of the Netherlands,” in *Proc. LREC-COLING (11th Workshop on the Representation and Processing of Sign Languages)*, pp. 273–281, 2024.
- [36] O. Ranum, J. Hadfield, and R. Bowden, “Spatial state collapse: Quantifying agreement failures under window truncation in streaming sign language translation,” *arXiv preprint arXiv:2606.08056*, 2026.
- [37] J.-P. Ryckaert, G. Ciccotti, and H. J. Berendsen, “Numerical integration of the cartesian equations of motion of a system with constraints: molecular dynamics of n-alkanes,” *Journal of Computational Physics*, vol. 23, no. 3, pp. 327–341, 1977.
- [38] Z. S. Sehyr, N. Caselli, A. M. Cohen-Goldberg, and K. Emmorey, “The ASL-LEX 2.0 Project: A visual and lexical database of American Sign Language,” *The Journal of Deaf Studies and Deaf Education*, vol. 26, no. 2, pp. 263–277, 2021.
- [39] H. Zhou et al., “SignBind: Aligning multimodal representations for sign language understanding with large language models,” in *Proc. ICLR*, 2025.
- [40] Y. Zhang, Y. Shen, et al., “Sign language translation using spatial configuration and motion dynamics with LLMs,” *arXiv preprint arXiv:2408.10593*, 2024.
- [41] R. Stiefelhagen and M. Finke, “Tracking focus of attention for human-robot interaction via eye gaze and pointing gestures,” in *Proc. IEEE/RSJ IROS*, pp. 1041–1046, 2002.
- [42] W. C. Stokoe, “Sign language structure: An outline of the visual communication systems of the American Deaf,” *Studies in Linguistics: Occasional Papers*, no. 8, 1960.
- [43] T. Supalla, “The classifier system of American Sign Language,” in *Noun Classification and Categorization*, C. Craig, Ed., John Benjamins, pp. 181–214, 1986.
- [44] G. Tevet, S. Raab, B. Gordon, Y. Shafir, D. Cohen-Or, and A. H. Bermano, “Human motion diffusion model (MDM),” in *Proc. ICLR*, 2023.
- [45] X. Liang et al., “ViPo-MLLM: Visual-pose multimodal large language model for continuous sign language translation,” in *Proc. IEEE/CVF CVPR*, 2026.
- [46] G. Xiao, Y. Tian, B. Chen, S. Han, and M. Lewis, “Efficient streaming language models with attention sinks,” in *Proc. ICLR*, 2024.
- [47] P. Xie, Q. Zhang, T. Peng, H. Tang, Y. Du, and Z. Li, “G2P-DDM: Generating sign pose sequence from gloss sequence with discrete diffusion model,” in *Proc. AAAI*, vol. 38, no. 6, pp. 6234–6242, 2024.
- [48] K. Yin and J. Read, “Better sign language translation with STMC-Transformer,” in *Proc. COLING*, pp. 5975–5989, 2020.
- [49] Y. Yuan, J. Song, U. Iqbal, A. Vahdat, and J. Kautz, “PhysDiff: Physics-guided human motion diffusion model,” in *Proc. IEEE/CVF ICCV*, pp. 11010–11021, 2023.
- [50] Y. Zhang, J. Lin, A. Zeng, G. Wu, S. Lu, Y. Fu, Y. Cai, R. Zhang, H. Wang, and L. Zhang, “Motion-X++: A large-scale multimodal 3D whole-body human motion dataset,” *arXiv preprint arXiv:2501.05098*, 2025.
- [51] G. Pavlakos, D. Shan, I. Radosavovic, A. Kanazawa, D. Fouhey, and J. Malik, “Reconstructing hands in 3D with transformers (HaMeR),” in *Proc. IEEE/CVF CVPR*, pp. 9826–9836, 2024.
- [52] J.-Y. Zhu, T. Park, P. Isola, and A. A. Efros, “Unpaired image-to-image translation using cycle-consistent adversarial networks,” in *Proc. IEEE/CVF ICCV*, pp. 2223–2232, 2017.